

Medical Insurance Cost Personal, Linear Machine Learning Model

August 17, 2023

```
[51]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib.pyplot as plt
%matplotlib inline
import warnings
warnings.filterwarnings("ignore")
```

Columns

age: age of primary beneficiary

sex: insurance contractor gender, female, male

bmi: Body mass index, providing an understanding of body, weights that are relatively high or low relative to height, objective index of body weight (kg / m^2) using the ratio of height to weight, ideally 18.5 to 24.9

children: Number of children covered by health insurance / Number of dependents

smoker: Smoking

region: the beneficiary's residential area in the US, northeast, southeast, southwest, northwest.

charges: Individual medical costs billed by health insurance

```
[132]: #To load the dataset
df = pd.read_csv(r"C:\Users\Aaditya Adyot\Downloads\Medical Cost Personal_
↳Datasets\insurance.csv")
```

```
[53]: # to see top 5 entries
df.head(5)
```

	age	sex	bmi	children	smoker	region	charges
0	19	female	27.900	0	yes	southwest	16884.92400
1	18	male	33.770	1	no	southeast	1725.55230
2	28	male	33.000	3	no	southeast	4449.46200
3	33	male	22.705	0	no	northwest	21984.47061
4	32	male	28.880	0	no	northwest	3866.85520

```
[54]: # to see last 5 entries
df.tail(5)
```

```
[54]:
```

	age	sex	bmi	children	smoker	region	charges
1333	50	male	30.97	3	no	northwest	10600.5483
1334	18	female	31.92	0	no	northeast	2205.9808
1335	18	female	36.85	0	no	southeast	1629.8335
1336	21	female	25.80	0	no	southwest	2007.9450
1337	61	female	29.07	0	yes	northwest	29141.3603

```
[55]: #to see the random sample in the dataset
df.sample(5)
```

```
[55]:
```

	age	sex	bmi	children	smoker	region	charges
1187	62	female	32.680	0	no	northwest	13844.79720
575	58	female	27.170	0	no	northwest	12222.89830
867	57	male	43.700	1	no	southwest	11576.13000
626	36	male	28.880	3	no	northeast	6748.59120
56	58	female	31.825	2	no	northeast	13607.36875

```
[56]: #to check the info of the dataset
df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 7 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   object
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   object
5   region      1338 non-null   object
6   charges     1338 non-null   float64
dtypes: float64(2), int64(2), object(3)
memory usage: 73.3+ KB
```

0.0.1 Observation: dtypes: float64(2), int64(2), object(3)

```
[57]: # to check the null values in the features
df.isna().sum()
```

```
[57]: age         0
sex         0
bmi         0
children    0
smoker      0
region      0
charges     0
dtype: int64
```

```
[58]: df[['smoker', 'region']]
```

```
[58]:      smoker    region
0         yes southwest
1         no  southeast
2         no  southeast
3         no northwest
4         no northwest
...
1333      no  northwest
1334      no  northeast
1335      no  southeast
1336      no southwest
1337      yes northwest

[1338 rows x 2 columns]
```

```
[59]: # to check the unique values in these feature
df['smoker'].unique()
```

```
[59]: array(['yes', 'no'], dtype=object)
```

```
[60]: df['region'].unique()
```

```
[60]: array(['southwest', 'southeast', 'northwest', 'northeast'], dtype=object)
```

```
[61]: ## for smoker and region, we can perform encoding
```

```
[62]: df['smoker'] = df['smoker'].map({'yes': 1, "no":0})
```

```
[66]: df= pd.concat([df, pd.get_dummies(df['region'])], axis =1)
```

```
[68]: df['sex'].unique()
```

```
[68]: array(['female', 'male'], dtype=object)
```

```
[71]: #use label encoder

from sklearn.preprocessing import LabelEncoder
encoder = LabelEncoder()
df['sex']=encoder.fit_transform(df['sex'])
```

```
[73]: df.drop('region', axis=1, inplace= True)
```

```
[74]: df
```

```
[74]:      age  sex    bmi  children  smoker    charges  northeast  northwest \
0      19    0  27.900         0        1  16884.92400         0         0
```

1	18	1	33.770	1	0	1725.55230	0	0
2	28	1	33.000	3	0	4449.46200	0	0
3	33	1	22.705	0	0	21984.47061	0	1
4	32	1	28.880	0	0	3866.85520	0	1
...
1333	50	1	30.970	3	0	10600.54830	0	1
1334	18	0	31.920	0	0	2205.98080	1	0
1335	18	0	36.850	0	0	1629.83350	0	0
1336	21	0	25.800	0	0	2007.94500	0	0
1337	61	0	29.070	0	1	29141.36030	0	1

	southeast	southwest
0	0	1
1	1	0
2	1	0
3	0	0
4	0	0
...
1333	0	0
1334	0	0
1335	1	0
1336	0	1
1337	0	0

[1338 rows x 10 columns]

```
[75]: df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1338 entries, 0 to 1337
Data columns (total 10 columns):
#   Column      Non-Null Count  Dtype
---  -
0   age         1338 non-null   int64
1   sex         1338 non-null   int32
2   bmi         1338 non-null   float64
3   children    1338 non-null   int64
4   smoker      1338 non-null   int64
5   charges     1338 non-null   float64
6   northeast   1338 non-null   uint8
7   northwest   1338 non-null   uint8
8   southeast   1338 non-null   uint8
9   southwest   1338 non-null   uint8
dtypes: float64(2), int32(1), int64(3), uint8(4)
memory usage: 62.8 KB
```

```
[76]: df.describe().T
```

```
[76]:
```

	count	mean	std	min	25% \
age	1338.0	39.207025	14.049960	18.0000	27.00000
sex	1338.0	0.505232	0.500160	0.0000	0.00000
bmi	1338.0	30.663397	6.098187	15.9600	26.29625
children	1338.0	1.094918	1.205493	0.0000	0.00000
smoker	1338.0	0.204783	0.403694	0.0000	0.00000
charges	1338.0	13270.422265	12110.011237	1121.8739	4740.28715
northeast	1338.0	0.242152	0.428546	0.0000	0.00000
northwest	1338.0	0.242900	0.428995	0.0000	0.00000
southeast	1338.0	0.272048	0.445181	0.0000	0.00000
southwest	1338.0	0.242900	0.428995	0.0000	0.00000

	50%	75%	max
age	39.000	51.000000	64.00000
sex	1.000	1.000000	1.00000
bmi	30.400	34.693750	53.13000
children	1.000	2.000000	5.00000
smoker	0.000	0.000000	1.00000
charges	9382.033	16639.912515	63770.42801
northeast	0.000	0.000000	1.00000
northwest	0.000	0.000000	1.00000
southeast	0.000	1.000000	1.00000
southwest	0.000	0.000000	1.00000

1 Detecting outliers

1. By using standard deviation method

```
[116]: index_list = []
feature = ['age', 'bmi', 'charges']
for i in feature:
    mean= df[i].mean()
    std = df[i].std()
    cutoff = std * 3
    index = df[(df[i] < (mean - cutoff)) | (df[i] > (mean + cutoff))].index
    index_list.extend(index)
```

```
[123]: index_list = sorted(set(index_list))
print(f'The no. outlier detected through Standard deviation method is, \
↳ {len(index_list)}')
print(f'The index of the dataset containing outliers are, {index_list}')
```

The no. outlier detected through Standard deviation method is, 11

The index of the dataset containing outliers are, [34, 116, 543, 577, 819, 847, 1047, 1146, 1230, 1300, 1317]

2. By using IQR

```
[125]: index_list = []
feature = ['age', 'bmi', 'charges']
for i in feature:
    lower_fence = df[i].quantile(.25)
    higher_fence= df[i].quantile(.75)
    IQR = higher_fence-lower_fence
    q1 = lower_fence - 1.5*IQR
    q2 = higher_fence + 1.5*IQR
    index = df[(df[i] < q1) | (df[i]> q2)].index
    index_list.extend(index)
```



```
[126]: index_list = sorted(set(index_list))
print(f'The no. outlier detected through Standard deviation method is, {len(index_list)}')
print(f'The index of the dataset containing outliers are, {index_list}')
```

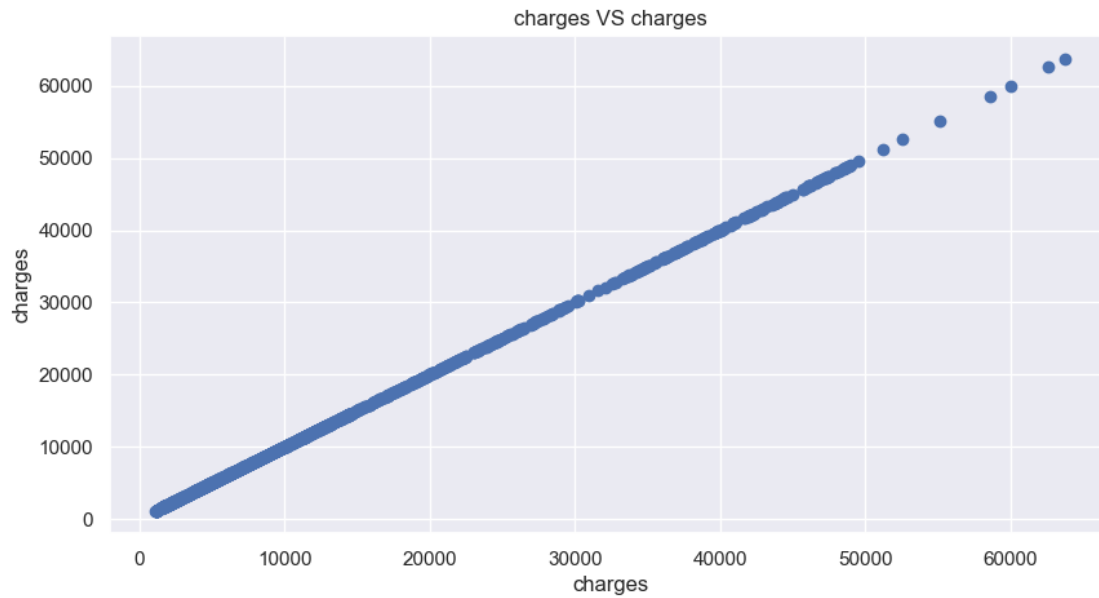
The no. outlier detected through Standard deviation method is, 145
The index of the dataset containing outliers are, [14, 19, 23, 29, 30, 34, 38, 39, 49, 53, 55, 82, 84, 86, 94, 109, 116, 123, 146, 158, 161, 175, 185, 203, 223, 240, 242, 251, 252, 254, 256, 263, 265, 271, 281, 286, 288, 292, 298, 312, 314, 322, 327, 328, 330, 338, 373, 377, 381, 401, 420, 421, 422, 441, 476, 488, 500, 524, 530, 543, 549, 558, 569, 577, 587, 609, 615, 621, 623, 629, 665, 667, 668, 674, 677, 682, 689, 697, 706, 725, 736, 738, 739, 742, 759, 803, 819, 826, 828, 842, 845, 847, 850, 852, 856, 860, 883, 893, 901, 917, 947, 951, 953, 956, 958, 1012, 1021, 1022, 1031, 1036, 1037, 1047, 1049, 1062, 1070, 1078, 1088, 1090, 1096, 1111, 1117, 1118, 1122, 1124, 1139, 1146, 1152, 1156, 1186, 1206, 1207, 1218, 1230, 1240, 1241, 1249, 1284, 1288, 1291, 1300, 1301, 1303, 1313, 1317, 1323]

3. using Using scatter for numerical feature and boxplot for categorical features

```
[127]: numerical_features = ['age', 'bmi', 'charges']
```

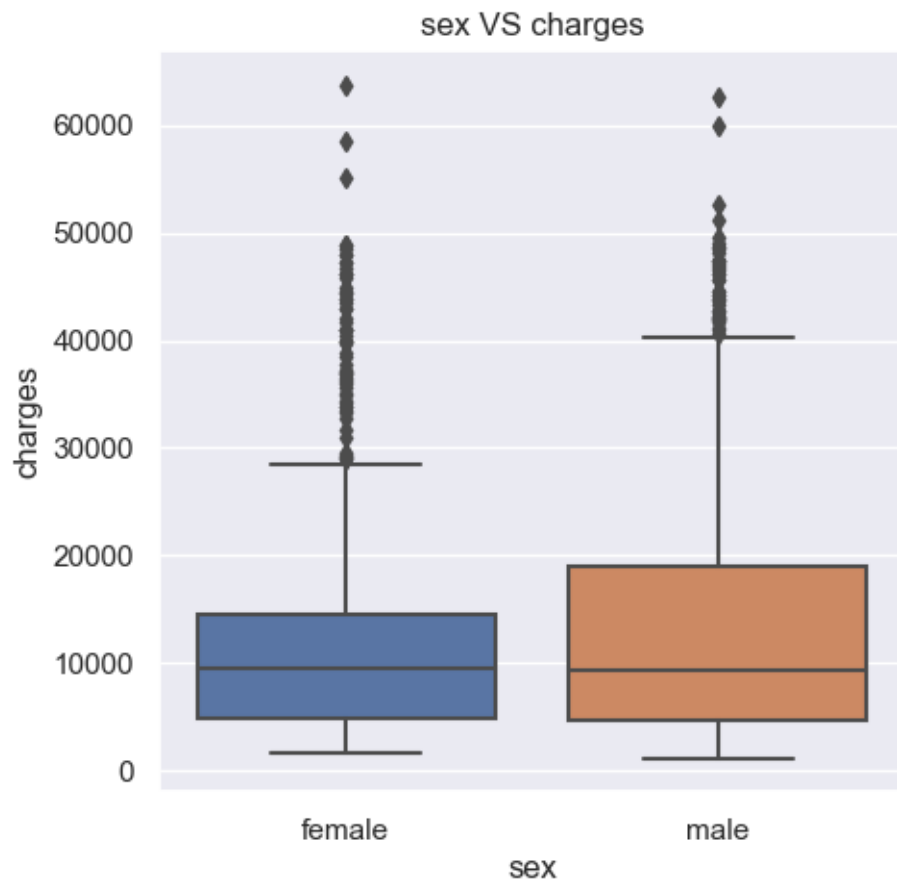
```
[129]: # numericals
for n in numerical_features :
    plt.figure(figsize=(10,5))
    plt.scatter(df[n], df["charges"])
    plt.title(f"{n} VS charges")
    plt.xlabel(f"{n}")
    plt.ylabel("charges")
    plt.show()
```

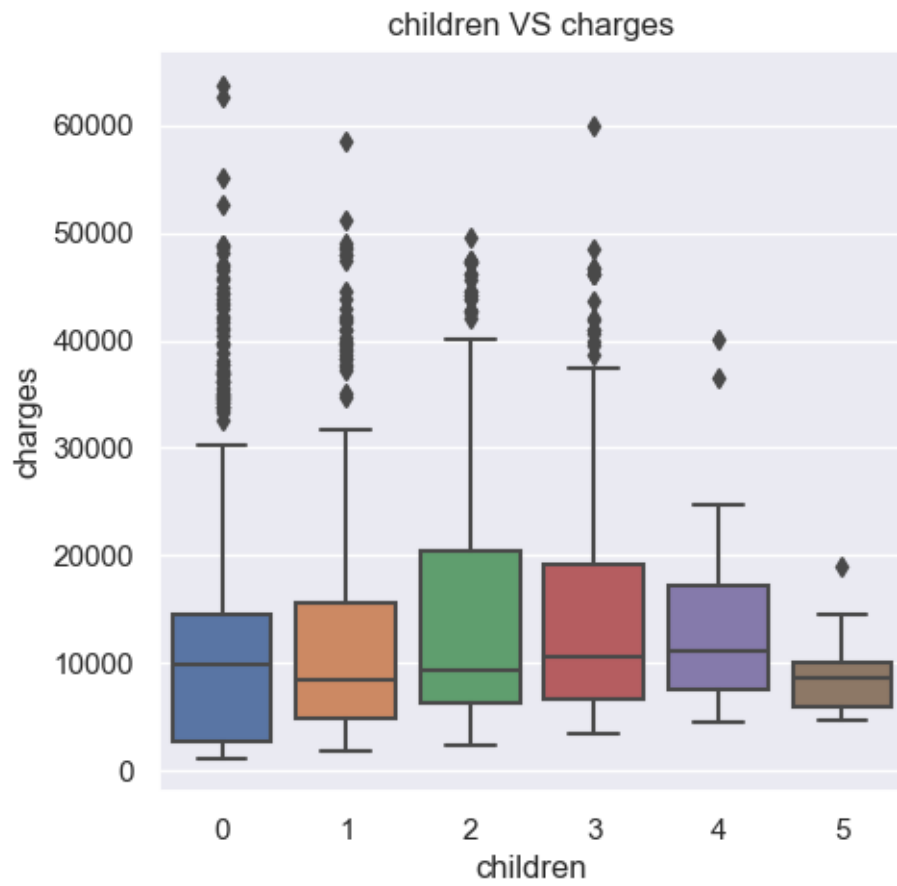


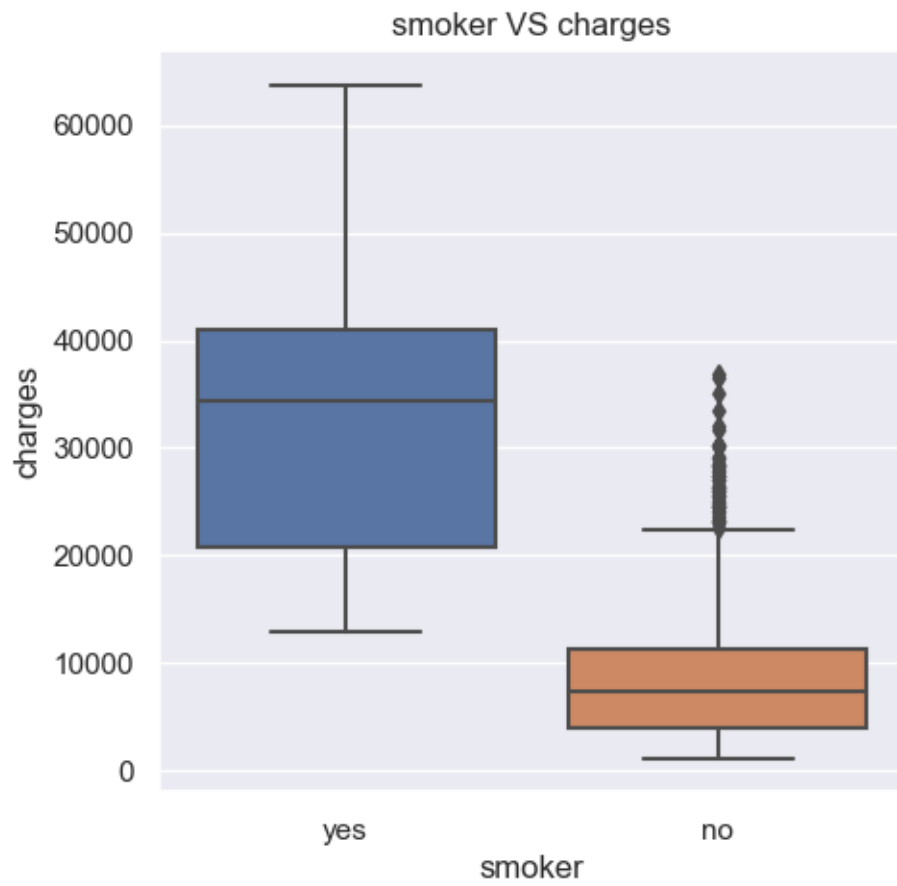


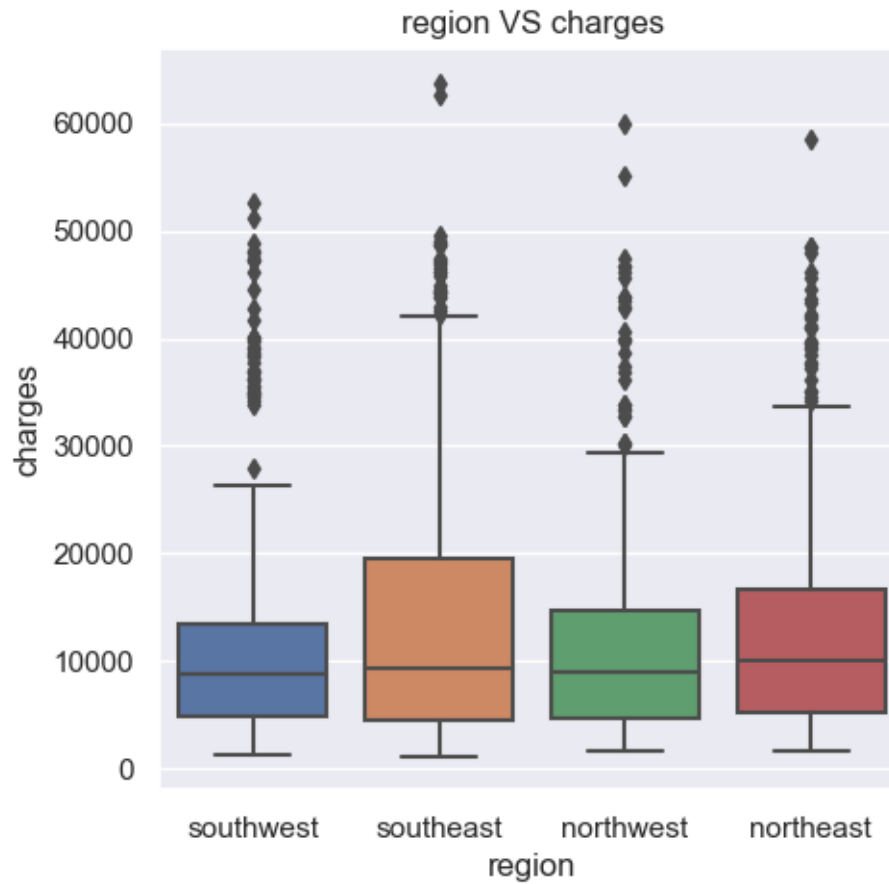
```
[130]: categorical_features = ["sex", "children", "smoker", "region"]
```

```
[135]: #categoricals
for n in categorical_features :
    plt.figure(figsize=(5,5))
    sns.boxplot(x=n, y="charges", data=df)
    plt.title(f"{n} VS charges")
    plt.show()
```








There are some outliers in our dataset. But it doesn't mean that they are false values or I should remove them. Also the number of samples are low.


```
[20]: sns.displot(df['charges'], kind='kde')
```

```
[20]: <seaborn.axisgrid.FacetGrid at 0x14c1c0e6920>
```




```
[22]: sns.boxplot(df['charges'], orient='h')
```

```
[22]: <Axes: >
```




```
[80]: sns.boxplot(df['age'], orient='h')
```

```
[80]: <Axes: >
```




```
[82]: sns.boxplot(df['bmi'], orient='h')
```

```
[82]: <Axes: >
```



```
[89]: sns.barplot(x='sex', y='smoker', data=df)
```

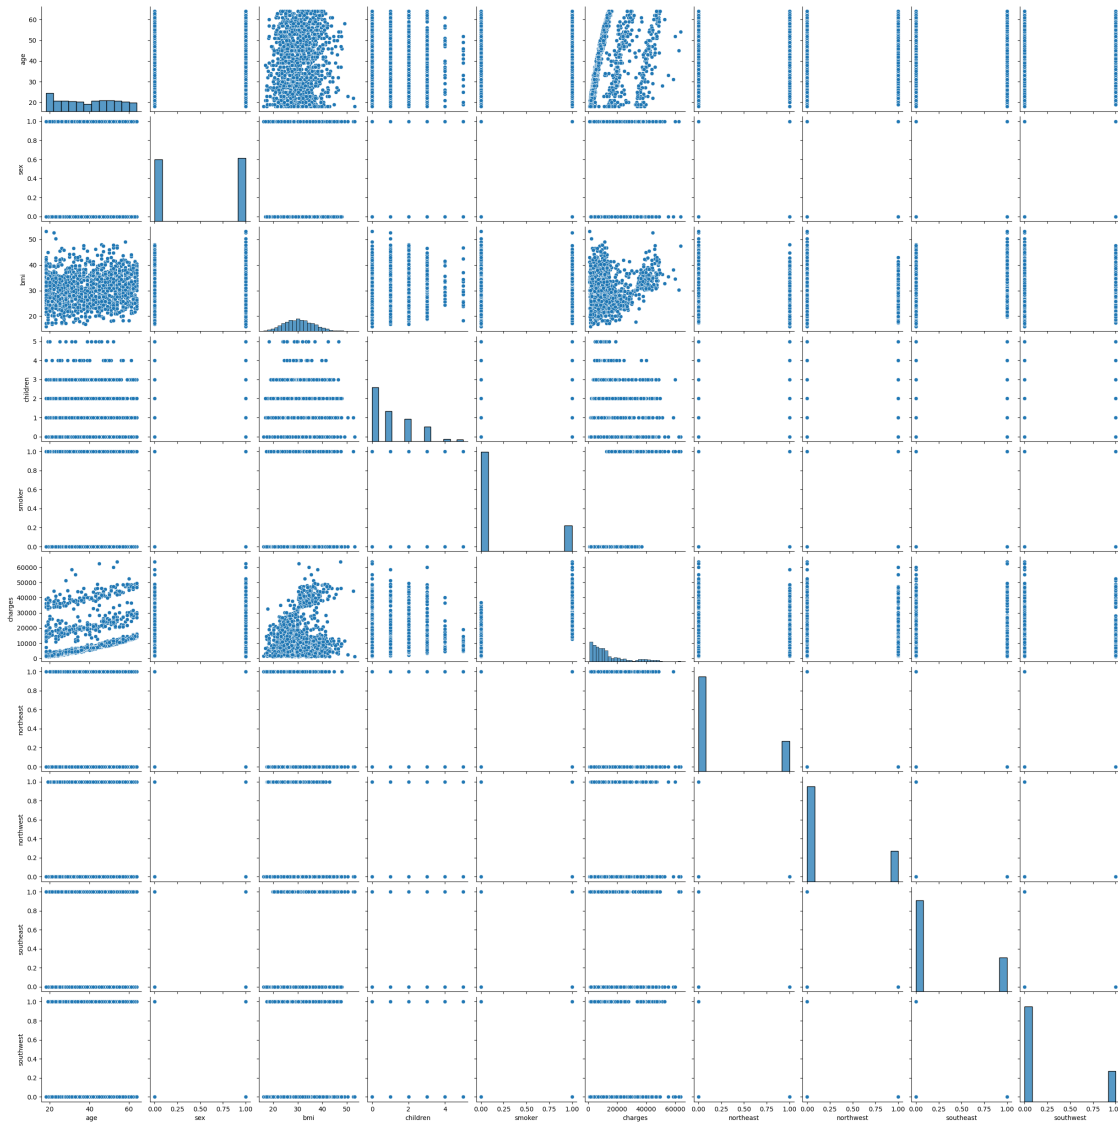
```
[89]: <Axes: xlabel='sex', ylabel='smoker'>
```

Observation: Male sex smokes more than the female sex.

```
[90]: sns.pairplot(df)
```

```
[90]: <seaborn.axisgrid.PairGrid at 0x14c271f3760>
```



```
[93]: df.corr()
```


```
[93]:
```

	age	sex	bmi	children	smoker	charges	\
age	1.000000	-0.020856	0.109272	0.042469	-0.025019	0.299008	
sex	-0.020856	1.000000	0.046371	0.017163	0.076185	0.057292	
bmi	0.109272	0.046371	1.000000	0.012759	0.003750	0.198341	
children	0.042469	0.017163	0.012759	1.000000	0.007673	0.067998	
smoker	-0.025019	0.076185	0.003750	0.007673	1.000000	0.787251	
charges	0.299008	0.057292	0.198341	0.067998	0.787251	1.000000	
northeast	0.002475	-0.002425	-0.138156	-0.022808	0.002811	0.006349	
northwest	-0.000407	-0.011156	-0.135996	0.024806	-0.036945	-0.039905	
southeast	-0.011642	0.017117	0.270025	-0.023066	0.068498	0.073982	
southwest	0.010016	-0.004184	-0.006205	0.021914	-0.036945	-0.043210	

	northeast	northwest	southeast	southwest
age	0.002475	-0.000407	-0.011642	0.010016
sex	-0.002425	-0.011156	0.017117	-0.004184
bmi	-0.138156	-0.135996	0.270025	-0.006205
children	-0.022808	0.024806	-0.023066	0.021914
smoker	0.002811	-0.036945	0.068498	-0.036945
charges	0.006349	-0.039905	0.073982	-0.043210
northeast	1.000000	-0.320177	-0.345561	-0.320177
northwest	-0.320177	1.000000	-0.346265	-0.320829
southeast	-0.345561	-0.346265	1.000000	-0.346265
southwest	-0.320177	-0.320829	-0.346265	1.000000

```
[92]: sns.set(rc={"figure.figsize": (15,10)})
sns.heatmap(df.corr(), annot=True)
```

```
[92]: <Axes: >
```



```
[94]: # Model building
X = df.drop('charges', axis =1)
y= df['charges']
```

```
[95]: from sklearn.model_selection import train_test_split
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42,
↪test_size=0.33)
```

```
[ ]: #Standardising the dataset
```

```
[96]: from sklearn.preprocessing import StandardScaler
scaler = StandardScaler()
X_train= scaler.fit_transform(X_train)
X_test= scaler.transform(X_test)
```

2 Applying the Liner regression

```
[97]: from sklearn.linear_model import LinearRegression
regression = LinearRegression()
regression.fit(X_train, y_train)
pred_data= regression.predict(X_test)
```

```
[98]: plt.scatter(y_test, pred_data)
plt.xlabel("actual data")
plt.ylabel("predicted data")
```


```
[98]: Text(0, 0.5, 'predicted data')
```



```
[99]: residuals = y_test - pred_data
```

```
[100]: sns.displot(residuals, kind= 'kde')
```

```
[100]: <seaborn.axisgrid.FacetGrid at 0x14c337e37f0>
```



```
[101]: from sklearn.metrics import mean_absolute_error, mean_squared_error
print(mean_squared_error(y_test, pred_data))
print(mean_absolute_error(y_test, pred_data))
print(np.sqrt(mean_squared_error(y_test, pred_data)))
```

```
35090225.72562568
```

```
4193.463021932157
```

```
5923.700340633857
```

```
[102]: from sklearn.metrics import r2_score
score = r2_score(y_test, pred_data)
score
```

```
[102]: 0.7605492639270064
```

3 Using Ridge (L2 regularisation)

```
[103]: from sklearn.linear_model import Ridge  
ridge = Ridge()  
ridge.fit(X_train, y_train)  
pred_data= ridge.predict(X_test)
```


```
[105]: plt.scatter(y_test, pred_data)  
plt.xlabel("Actual data output")  
plt.ylabel("Predicted data output")
```

```
[105]: Text(0, 0.5, 'Predicted data output')
```



```
[107]: residuals= y_test - pred_data  
sns.displot(residuals, kind= 'kde')
```

```
[107]: <seaborn.axisgrid.FacetGrid at 0x14c3325d8d0>
```



```
[108]: from sklearn.metrics import mean_absolute_error, mean_squared_error
print(mean_squared_error(y_test, pred_data))
print(mean_absolute_error(y_test, pred_data))
print(np.sqrt(mean_squared_error(y_test, pred_data)))
```

```
35091586.871670134
4194.950790847867
5923.815229366133
```

```
[109]: from sklearn.metrics import r2_score
score = r2_score(y_test, pred_data)
score
```


```
[109]: 0.7605399756589628
```

4 using Lasso (L1 Regularisation)

```
[110]: from sklearn.linear_model import Lasso
lasso = Lasso()
lasso.fit(X_train, y_train)
pred_data= lasso.predict(X_test)
```


```
[111]: plt.scatter(y_test, pred_data)
plt.xlabel("Actual data output")
plt.ylabel("Predicted data output")
```

```
[111]: Text(0, 0.5, 'Predicted data output')
```



```
[112]: residuals= y_test - pred_data
sns.displot(residuals, kind= 'kde')
```

```
[112]: <seaborn.axisgrid.FacetGrid at 0x14c33b671f0>
```

```
[113]: from sklearn.metrics import mean_absolute_error, mean_squared_error
print(mean_squared_error(y_test, pred_data))
print(mean_absolute_error(y_test, pred_data))
print(np.sqrt(mean_squared_error(y_test, pred_data)))
```

```
35090559.44369504
4193.3943816846895
5923.72850860799
```

```
[114]: from sklearn.metrics import r2_score
score = r2_score(y_test, pred_data)
score
```


```
[114]: 0.7605469866821127
```

5 ElasticNet Linear Model

```
[136]: from sklearn.linear_model import ElasticNet
      EN = ElasticNet()
      EN.fit(X_train, y_train)
      pred_data = EN.predict(X_test)
```

```
[137]: plt.scatter(y_test, pred_data)
      plt.xlabel("Actual test data")
      plt.ylabel("predicted test data")
```


```
[137]: Text(0, 0.5, 'predicted test data')
```



```
[138]: residual = y_test - pred_data
```

```
[139]: sns.displot(residual, kind = 'kde')
```

```
[139]: <seaborn.axisgrid.FacetGrid at 0x14c33910850>
```



```
[140]: from sklearn.metrics import mean_absolute_error, mean_squared_error
print(mean_squared_error(y_test, pred_data))
print(mean_absolute_error(y_test, pred_data))
print(np.sqrt(mean_squared_error(y_test, pred_data)))
```

```
47985463.21529332
5190.321365611101
6927.154048762978
```

```
[141]: from sklearn.metrics import r2_score
score = r2_score(y_test, pred_data)
score
```

```
[141]: 0.6725539876104439
```

```
[146]: # adjusted r2 score
1-(1-score)*(len(y_test)-1)/ (len(y_test)-X_test.shape[1]-1)
```

```
[146]: 0.6657321956856616
```

[]:

[]:

[]: